6.5620 (6.875), Fall 2022

Homework # 1 **Due: 21 September 2022, 11:59:59pm ET**

- **Typsetting:** You are encouraged to use LaTeX to typeset your solutions. You can use the following template.
- **Submissions:** Solutions should be submitted to Gradescope.
- **Reference your sources:** If you use material outside the class, please reference your sources (including papers, websites, wikipedia).
- **Acknowledge your collaborators:** Collaboration is permitted and encouraged in small groups of at most three. You must write up your solutions entirely on your own and acknowledge your collaborators.

Problems:

1. *(5 points)* **Working with negligible functions.** Recall that a non-negative function $\nu : \mathbb{N} \to \mathbb{R}^+$ is *negligible* if it decreases faster than the inverse of any polynomial (otherwise, we say that ν is *non-negligible*). More precisely, for all $c > 0$, there exists a constant $N > 0$ such that for all $n \geq N$, $\nu(n) < n^{-c}$.

   State whether each of the following functions is negligible or non-negligible, and prove your assertion. For all of the problems below, we take the base of the logarithm to be 2.

   (a) $\nu(n) = 1/2^{100 \log n}$.

   (b) $\nu(n) = 1/2^{100 \log n \cdot \log \log \log n}$.

   (c) $\nu(n) = p(n) \cdot \mu(n)$, where $p(n) = O(n^k)$ for some constant k, and $\mu(n)$ is a negligible function. Either prove that ν is always negligible, or come up with a counter-example.

   (d) $\nu(n) = (\mu(n))^{\frac{1}{p(n)}}$, where $p(n)$ and $\mu(n)$ are as defined in (c). Either prove that ν is always negligible, or come up with a counter-example.

   (e) $\nu(n) = 1/2^{2^{\log^* n}}$, where $\log^* n$ is the number of times the logarithm function must be iteratively applied to n before the result is less than or equal to 1. More concretely,

   $$\log^* n := \begin{cases} 0 & \text{if } n \leq 1 \\ 1 + \log^*(\log n) & \text{if } n > 1. \end{cases}$$

   You may use the fact that $\log^* n \leq n$ without proof. (Hint: How does $\log^* n$ compare to $\log \log n$?)

2. *(11 points)* **Statistical and computational indistinguishability.** We think of distributions X, Y on a (finite) set Ω as functions $X, Y : \Omega \to [0,1]$ such that for $\sum_{\omega\in\Omega} X(w) = \sum_{\omega\in\Omega} Y(\omega) = 1$. The statistical distance (also known as variational or L_1 distance) between X and Y is defined as

$$\Delta(X,Y) := \frac{1}{2}\sum_{\omega\in\Omega} |X(\omega) - Y(\omega)| \ .$$

   (a) *(3 points)* Show that the following is an equivalent definition:

$$\Delta(X,Y) := \sup_{A\subseteq\Omega} |X(A) - Y(A)| \ ,$$

   where $X(A)$ is shorthand for $\sum_{\omega\in A} X(\omega)$.

   (b) *(3 points)* Let D_0 and D_1 be two distributions over Ω. Suppose that we play the following game with an algorithm $\mathcal{A}$. First, we pick at random a bit $b \leftarrow \{0,1\}$ and then we pick $x \leftarrow D_b$ and we give x to $\mathcal{A}$. Finally, $\mathcal{A}$ returns a bit. It wins if the bit returned is equal to b. Show that the highest success probability in this game is exactly $\frac{1}{2} + \frac{1}{2}\Delta(D_0, D_1)$.

   (c) *(1 point)* Give the definition of computational indistinguishability using similar language as in the previous question. (This part should not take more than 3-5 sentences.)

   (d) *(4 points)* For a probability distribution D over Ω and positive integer m, let D^m denote the *product distribution* over Ω^m, obtained by drawing a tuple of m independent samples from D. Let $\mathcal{X} = \{X_n\}_n$ and $\mathcal{Y} = \{Y_n\}_n$ be ensembles of distributions that are efficiently sampleable (in PPT), and let $m(n) = \mathrm{poly}(n)$ be some fixed polynomial. Prove that if $\mathcal{X}$ and $\mathcal{Y}$ are computationally indistinguishable, or, in symbols, $\mathcal{X} \overset{c}{\approx} \mathcal{Y}$, then $\{X_n^{m(n)}\} \overset{c}{\approx} \{Y_n^{m(n)}\}$. (Where do you use that X_n, Y_n are efficiently sampleable?) This shows that if two efficiently sampleable distributions are computationally indistinguishable given one sample, then they are also computationally indistinguishable given polynomially many samples.

3. *(8 points)* **PRG or not?** Let $G : \{0,1\}^{2n} \to \{0,1\}^{\ell}$ be a pseudorandom generator, where $\ell \geq 2n+1$. In each of the following, say whether G_c is necessarily a pseudorandom generator. If yes, give a proof. Otherwise, show a counterexample. Your counterexamples must rely only on the existence of pseudorandom generators.

   (a) *(2 points)* Consider $G_0 : \{0,1\}^{2n} \to \{0,1\}^{\ell}$, where $G_0(s) := G(\overline{s})$. Here, $\overline{s}$ is the bit-wise complement of s.

   (b) *(3 points)* Consider $G_1 : \{0,1\}^{n} \to \{0,1\}^{\ell}$, where $G_1(s) := G(0^n||s)$.

   (c) *(3 points)* Consider $G_2 : \{0,1\}^{2n} \to \{0,1\}^{2\ell}$, where $G_2(s) := G(s)||G(s+1 \mod 2^{2n})$.

4. *(3 points)* **A PRG from a PRF.** Prove that, if F is a length preserving pseudorandom function, then $G(s) \stackrel{\text{def}}{=} F_s(\langle 1\rangle)\|F_s(\langle 2\rangle)\|\ldots\|F_s(\langle \ell\rangle)$, where $\langle i\rangle$ is the n-bit binary representation of i, is a pseudorandom generator with expands ℓ bits to $\ell \cdot n$ bits.

5. *(8 points)* **Malicious Mind-Reading.**

   Veronica claims to Lucy, her classmate in 6.5620, that she can read minds, but Lucy does not believe it one bit, and is willing to bet \$100 that it's all baloney. They decide to play a game where Veronica can prove to Lucy that she is a bonafide mind-reader. If Veronica is lying, she should not be able to win in the game; and if Veronica indeed possesses this supernatural power, Lucy should grant victory (and \$100) to Veronica at the end of the game (no matter how much she hates to lose). They decide to play the following game.

   (a) Veronica (we'll call her V) sends Lucy a string v (of a certain length that they decided on), chosen at random.

   (b) Lucy (we'll call her L) chooses a random bit $\sigma \in \{0, 1\}$, and Veronica has to guess what this bit is. To do this, Lucy sends

   $$\ell \leftarrow L(\sigma, v; r)$$

   to Veronica, where r is Lucy's private random coins that only she knows.

   (c) Now, Veronica reads Lucy's mind and guesses what σ is.

   (d) Finally, Lucy "unlocks" her bit σ by sending σ and r to Veronica. Veronica verifies that ℓ is indeed $L(\sigma, v; r)$.

   Veronica wins if her guess of σ is correct.

   (a) *(1 point)* Veronica should not gain any information about Lucy's bit from viewing the lock (i.e. after step 2 of the game). In other words, a malicious (but computationally *bounded*) Veronica V^* should not be able to learn anything about the honest L's choice bit σ, no matter what initial message v^* she sent.

   Using the notion of indistinguishability, give a formal definition of this *concealing property* of L.

   (b) *(1 point)* Lucy should not be able to unlock her bit both ways, otherwise she can always get away with not paying Veronica \$100 even if Veronica is a mind-reader (do you see why?) To ensure this, the locking algorithm L has to be "unmodifiable". Give a formal definition of this property.

   (c) *(2 points)* Let G be any length-tripling function, i.e., one for which $|G(x)| = 3|x|$ for every $x \in \{0, 1\}^*$. Give an upper bound on the probability, over the choice of a random $3n$-bit string v, that there exist two inputs $x_1, x_2 \in \{0, 1\}^n$ such that $G(x_1) \oplus G(x_2) = v$.

(d) *(4 points)* Let G be a length-tripling PRG (which we have seen can be obtained from any PRG). Use G to construct a secure locking scheme (i.e. define the algorithms V and L), and prove that it is both concealing and unmodifiable according to your definitions.