# 6.5620 (6.875), Fall 2022

**Homework # 6**                                      **Due: Dec 7 2022, 11:59:59pm ET**

- **Typsetting:** You are encouraged to use LaTeX to typeset your solutions. You can use the following template.

- **Submissions:** Solutions should be submitted to Gradescope.

- **Reference your sources:** If you use material outside the class, please reference your sources (including papers, websites, wikipedia).

- **Acknowledge your collaborators:** Collaboration is permitted and encouraged in small groups of at most three. You must write up your solutions entirely on your own and acknowledge your collaborators.

---

**Problems:**

1. *(8 points)* **Solution-checking systems.**

   Clint and Mallory are pset partners for 6.875. They enjoy collaborating on psets a lot, but they often work alone a bit before discussion, to test their individual problem-solving skills. Also, they are usually offline on Messenger, so they communicate in the non-interactive setting.

   Lest they spoil the fun of problem-solving, they agree to use *zero-knowledge solution-checking systems* when they try to convince each other that they have already discovered some solution $w$ (i.e. a witness) to a problem $x \in L$ ($L$ is an NP language) in the problem set.

   A *solution-checking system* is a tuple of p.p.t. algorithms $\Pi = (\mathsf{Setup}, P, V)$ with the following usage.

   (i) The trusted TAs compute $(\mathsf{pp}, \mathsf{sp}) \leftarrow \mathsf{Setup}(1^n)$, send $\mathsf{pp}$ to Mallory, and send $(\mathsf{pp}, \mathsf{sp})$ to the Clint.

   (ii) Mallory discovers a witness $w$ certifying that $x \in L$, and wants to prove to Clint that she knows $w$.

   (iii) Mallory computes $\pi \leftarrow P(1^n, \mathsf{pp}, x, w)$ and sends $\pi$ to Clint.

   (iv) Clint accepts iff $V(1^n, \mathsf{sp}, x, \pi) = 1$. (We require that $V$ computes a deterministic predicate.)

Note that the output of Setup does not depend on the statement $x \in L$ (nor $w$), as Setup was run by the trusted TAs before the problem sets were even written up. Also note that the system is designed to be used on a per-verifier basis, as sp is only sent to Clint but not Mallory (nor anyone else in the class). If Clint wanted to prove to Mallory that he found a solution, they would have to switch roles.

Assume we already have a tuple of p.p.t algorithms $\Pi_0 = (P_0^1, P_0^2, V_0)$ with the following usage.

(i) An honest prover has an input $x$, and a witness $w$ certifying that $x \in L$.

(ii) The prover computes $\alpha \leftarrow P_0^1(1^n, x, w)$ and sends $\alpha$ to the verifier.

(iii) The verifier samples a random bit $\beta \leftarrow_R \{0, 1\}$ and sends $\beta$ to the prover.

(iv) The prover computes $\gamma \leftarrow P_0^2(1^n, x, w, \alpha, \beta)$ and sends it to the verifier.

(v) The verifier accepts iff $V_0(1^n, \alpha, \beta, \gamma) = 1$, where $V_0$ is a deterministic predicate.

As a reminder, $|\alpha|, |\gamma| \in \text{poly}(n)$. $\Pi_0$ satisfies **perfect completeness**, **soundness with error** $1/2$, and **computational honest-verifier zero-knowledge** with simulator $\mathcal{S}$ (i.e., there exists PPT $\mathcal{S}$ such that $\{(\tilde{\alpha}, \tilde{\beta}, \tilde{\gamma}) : (\tilde{\alpha}, \tilde{\beta}, \tilde{\gamma}) \leftarrow \mathcal{S}(1^n, x)\}$ is computationally indistinguishable from a transcript of a real execution between the honest verifier $V$ and the honest prover $P$).

Such a protocol for any NP language can be constructed if we try hard enough to construct a zero-knowledge proof system for the NP-complete problem Hamiltonian Path, using ideas similar to the GMW protocol for 3-coloring. Fortunately, our friend Manuel open-sourced and provided for free such a protocol $\Pi_0$, so we can just use it for free.

**Your task:** Transform $\Pi_0$ into a solution-checking system $\Pi$, using any IND-CPA secure public-key encryption scheme $\mathcal{E} = (\text{Gen}, \text{Enc}, \text{Dec})$, so that Clint and Mallory can show off to each other that they solved problems without spoiling the fun of problem-solving for the other person.

In particular, prove that your solution-checking system $\Pi = (\text{Setup}, P, V)$ satisfies *perfect completeness*, *soundness with negligibly small soundness error*, and *per-verifier zero-knowledge* with simulator $\mathcal{S}$:

(i) *Perfect completeness.* For every $x \in L$ with witness $w$,

$$\mathbf{Pr}[(\text{pp}, \text{sp}) \leftarrow \text{Setup}(1^n); \pi \leftarrow P(1^n, \text{pp}, x, w) : V(1^n, \text{pp}, \text{sp}, x, \pi) = 1] = 1.$$

(ii) *Soundness.* For every $x \notin L$ and any purported proof $\pi$,

$$\mathbf{Pr}[(\text{pp}, \text{sp}) \leftarrow \text{Setup}(1^n) : V(1^n, \text{pp}, \text{sp}, x, \pi) = 1] < \text{negl}(n).$$

(iii) *Per-verifier zero-knowledge.* For every $x \in L$ and any witness $w$ certifying that $x \in L$,

$$\mathcal{D}_1 \leftarrow \{(\widetilde{\mathsf{pp}}, \widetilde{\mathsf{sp}}, x, \tilde{\pi}) : (\widetilde{\mathsf{pp}}, \widetilde{\mathsf{sp}}, x, \tilde{\pi}) \leftarrow \mathcal{S}(1^n, x)\},$$

and

$$\mathcal{D}_2 \leftarrow \{(\mathsf{pp}, \mathsf{sp}, x, \pi) : (\mathsf{pp}, \mathsf{sp}) \leftarrow \mathsf{Setup}(1^n, x), \pi \leftarrow P(1^n, \mathsf{pp}, x, w)\}$$

are computationally indistinguishable.

You may assume for simplicity that $|x| = n$.

*Remark.* (Note this is not a supposed to be a hint.) To simplify presentation, all computational security notions are quantified over non-uniform p.p.t. adversaries in this problem.

*Remark.* (Note this is not a supposed to be a hint.) It might be interesting to ponder if your system is *reusable*, in the sense that when Clint and Mallory use the same setup parameters more than once, soundness and (per-verifier) zero-knowledge still hold. Or more generally, if you want to use the system in a bigger system, what might become problematic? (Generally such considerations are called *composability.*)

*Remark.* (Note this is not a supposed to be a hint.) For simplicity, we have only asked you to prove ordinary soundness in this problem, but in order to properly capture their application, Clint and Mallory should actually design a zero-knowledge proof system that is a *proof of knowledge.* Note that basic soundness only guarantees that false statements cannot be proven—and, if the TAs are to be trusted, then $x$ is always in $L$ anyway, so there is no point to having Mallory prove that $x \in L$ to Clint. On the other hand, the stronger property of *proof of knowledge soundness* (or *knowledge soundness*) requires that there exists a p.p.t. algorithm, called an *extractor*, that with rewinding oracle access to any potentially cheating prover $P^*$, can output a valid witness $w$ certifying that $x \in L$. The existence of the extractor can then be taken as proof that, if Mallory is accepted by Clint, she must have actually 'known' the witness $w$, and not merely the fact that $x \in L$.

In practice, many zero-knowledge proof systems are also proofs of knowledge (or can be fairly easily modified to capture the stronger soundness requirement). For simplicity, we do not ask you to prove the stronger soundness property here.

## Solution

Define $\Pi = (\mathsf{Setup}, P, V)$ as follows.

- $\mathsf{Setup}(1^n)$:

  (i) Sample $2n$ public-secret key pairs, $(\mathsf{pk}_i^b, \mathsf{sk}_i^b) \leftarrow \mathsf{Gen}(1^n)$ for $b \in \{0, 1\}, i \in [n]$.

  (ii) Sample $\beta \leftarrow_R \{0, 1\}^n$.

  (iii) Output $\mathsf{pp} \leftarrow (\mathsf{pk}_i^b)_{b \in \{0,1\}, i \in [n]}$, $\mathsf{sp} \leftarrow (\beta_i, \mathsf{sk}_i^{\beta_i})_{i \in [n]}$.

- $P(1^n, \mathsf{pp}, x, w)$:

  (i) Parse $(\mathsf{pk}_i^b)_{b \in \{0,1\}, i \in [n]} \leftarrow \mathsf{pp}$.

  (ii) Sample $n$ first-messages $\alpha_i \leftarrow P_0^1(1^n, x, w)$ for $i \in [n]$.

  (iii) Compute $\gamma_i^b \leftarrow P_0^2(1^n, x, w, \alpha_i, b)$ for $b \in \{0, 1\}, i \in [n]$.

  (iv) Output $\pi \leftarrow (\alpha_i, \mathsf{Enc}(\mathsf{pk}_i^b, \gamma_i^b))_{b \in \{0,1\}, i \in [n]}$.

- $V(1^n, \mathsf{pp}, \mathsf{sp}, x, \pi)$:

  (i) Parse $(\beta_i, \mathsf{sk}_i^b)_{b \in \{0,1\}, i \in [n]} \leftarrow \mathsf{sp}$.

  (ii) Parse $(\alpha_i, c_i^b)_{b \in \{0,1\}, i \in [n]} \leftarrow \pi$.

  (iii) Compute $\gamma_i^b \leftarrow \mathsf{Dec}(\mathsf{sk}_i^b, c_i^b)$ for $b \in \{0, 1\}, i \in [n]$.

  (iv) Accept iff $V_0(1^n, \alpha_i, \beta_i, \gamma_i^{\beta_i}) = 1$ for all $i \in [b]$.

**Completeness** is immediate by the completeness of $\Pi_0$.

**Soundness** Note that $\Pi_0$ has soundness error $1/2$, this means there is at least one challenge $b_i \in \{0, 1\}$ that would be rejected by $V_0$, if $x \notin L$. Moreover, $\mathsf{pp}$ is independent from $\beta \in \{0, 1\}^n$, so Mallory would have no information about it, hence any (unbounded prover) would be caught with probability $1 - 2^{-n}$, as each of the $n$ invalid responses would be caught with probability $1/2$.

## Solution

**Per-verifier Zero-knowledge**   Assume $x \in L$.

- $\mathcal{D}_1 \leftarrow \mathcal{S}(1^n, x)$:

  (i) Sample $2n$ public-secret key pairs, $(\mathsf{pk}_i^b, \mathsf{sk}_i^b) \leftarrow \mathsf{Gen}(1^n)$ for $b \in \{0,1\}, i \in [n]$.

  (ii) Compute $(\tilde{\alpha}_i, \tilde{\beta}_i, \tilde{\gamma}_i) \leftarrow \mathcal{S}_0(1^n, x)$ for $i \in [n]$.

  (iii) Compute $\widetilde{\mathsf{pp}} \leftarrow (\mathsf{pk}_i^b)_{b \in \{0,1\}, i \in [n]}$, $\widetilde{\mathsf{sp}} \leftarrow (\tilde{\beta}_i, \mathsf{sk}_i^{\tilde{\beta}_i})_{i \in [n]}$.

  (iv) Compute $\tilde{c}_i^{\tilde{\beta}_i} \leftarrow \mathsf{Enc}(\mathsf{pk}_i^{\tilde{\beta}_i}, \tilde{\gamma}_i^{\tilde{\beta}_i})$, and $\tilde{c}_i^{(1-\tilde{\beta}_i)} \leftarrow \mathsf{Enc}(\mathsf{pk}_i^{1-\tilde{\beta}_i}, 0^{|\gamma|})$ for every $i \in [n]$.

  (v) Output $(\widetilde{\mathsf{pp}}, \widetilde{\mathsf{sp}}, x, \tilde{\pi} := (\tilde{\alpha}_i, \tilde{c}_i^b)_{b \in \{0,1\}, i \in [n]})$.

Assume there exists a (non-uniform) p.p.t. distinguisher $\mathcal{A}$ that distinguishes $(\mathcal{D}_1, \mathcal{D}_2)$ when the statement-witness pair is $(x, w)$.
Note that this reduction is non-uniform because $(x, w)$ is given as an advice. We define the following hybrids.

- $\mathcal{D}_1^j$, for each $j \in \{0, 1, \ldots, n\}$ :

  (i) Sample $2n$ public-secret key pairs, $(\mathsf{pk}_i^b, \mathsf{sk}_i^b) \leftarrow \mathsf{Gen}(1^n)$ for $b \in \{0,1\}, i \in [n]$.

  (ii) Compute $(\tilde{\alpha}_i, \tilde{\beta}_i, \tilde{\gamma}_i) \leftarrow \mathcal{S}_0(1^n, x)$ for $i \in \{1, \ldots, j\}$.

  (iii) Compute $(\tilde{\alpha}_i, \tilde{\beta}_i, \tilde{\gamma}_i) \leftarrow P_0^1(1^n, x, w)$ for $i \in \{j+1, \ldots, n\}$.

  (iv) Compute $\widetilde{\mathsf{pp}} \leftarrow (\mathsf{pk}_i^b)_{b \in \{0,1\}, i \in [n]}$, $\widetilde{\mathsf{sp}} \leftarrow (\tilde{\beta}_i, \mathsf{sk}_i^{\tilde{\beta}_i})_{i \in [n]}$.

  (v) Compute $\tilde{c}_i^{\tilde{\beta}_i} \leftarrow \mathsf{Enc}(\mathsf{pk}_i^{\tilde{\beta}_i}, \tilde{\gamma}_i^{\tilde{\beta}_i})$, and $\tilde{c}_i^{(1-\tilde{\beta}_i)} \leftarrow \mathsf{Enc}(\mathsf{pk}_i^{1-\tilde{\beta}_i}, 0^{|\gamma|})$ for every $i \in [n]$.

  (vi) Output $(\widetilde{\mathsf{pp}}, \widetilde{\mathsf{sp}}, x, \tilde{\pi} := (\tilde{\alpha}_i, \tilde{c}_i^b)_{b \in \{0,1\}, i \in [n]})$.

- $\mathcal{D}_2^j$, for each $j \in \{0, 1, \ldots, n\}$ :

  (i) Sample $2n$ public-secret key pairs, $(\mathsf{pk}_i^b, \mathsf{sk}_i^b) \leftarrow \mathsf{Gen}(1^n)$ for $b \in \{0,1\}, i \in [n]$.

  (ii) Compute $(\tilde{\alpha}_i, \tilde{\beta}_i, \tilde{\gamma}_i) \leftarrow P_0^1(1^n, x, w)$ for $i \in [n]$.

  (iii) Compute $\widetilde{\mathsf{pp}} \leftarrow (\mathsf{pk}_i^b)_{b \in \{0,1\}, i \in [n]}$, $\widetilde{\mathsf{sp}} \leftarrow (\tilde{\beta}_i, \mathsf{sk}_i^{\tilde{\beta}_i})_{i \in [n]}$.

  (iv) Compute $\tilde{c}_i^{\tilde{\beta}_i} \leftarrow \mathsf{Enc}(\mathsf{pk}_i^{\tilde{\beta}_i}, \tilde{\gamma}_i^{\tilde{\beta}_i})$,

  (v) Compute $\tilde{c}_i^{(1-\tilde{\beta}_i)} \leftarrow \mathsf{Enc}(\mathsf{pk}_i^{1-\tilde{\beta}_i}, 0^{|\gamma|})$ for every $i \in \{1, \ldots, j\}$.

  (vi) Compute $\tilde{c}_i^{(1-\tilde{\beta}_i)} \leftarrow \mathsf{Enc}(\mathsf{pk}_{i\,5}^{1-\tilde{\beta}_i}, \gamma_i^{1-\tilde{\beta}_i})$ for every $i \in \{j+1, \ldots, n\}$.

  (vii) Output $(\widetilde{\mathsf{pp}}, \widetilde{\mathsf{sp}}, x, \tilde{\pi} := (\tilde{\alpha}_i, \tilde{c}_i^b)_{b \in \{0,1\}, i \in [n]})$.

All the adjacent hybrids are indistinguishable, by the securities of $\mathcal{S}$ and IND-CPA security.

2. *(11 points)* **Two-Server PIR with Pre-Processing**

In this problem, we will consider private information retrieval. In this model, an array $x \in \{0,1\}^n$ is stored by one or more servers, and a user wants to a single bit of x, without the servers knowing which bit he wants. We will consider a two-server model of PIR. Each server has a copy of $x$ and will answer the users queries, but the servers **do not communicate** with each other. Consider the following two-server PIR protocols between user U and servers $S_0$ and $S_1$:

- The user chooses a random subset $T \subseteq \{1, 2, \ldots, n\}$, including each value independently with probability 1/2. (Note that $T$ can be represented as an $n$-bit string $w_T$ by letting the bit in position $j$ represent whether or not $j \in T$.)

- The user calculates the set $T \oplus i$, which is $T \cup \{i\}$ if $T$ does not contain $i$, and $T \setminus \{i\}$ if $T$ does not contain $i$.

- The user sends $w_T$ to the server $S_0$ and $w_{T \oplus i}$ to the server $S_1$.

- Given a set, each server sends back the XOR of all the bits of $x$ in that set.

- The user calculates the XOR of the two responses from $S_0$ and $S_1$.

(a) *(2 points)* Prove that the scheme is *information-theoretically private*; that is, for any two different values $i$ and $i'$, the view of each server $S_0$ and $S_1$ is perfectly indistinguishable, even if the servers are computationally unbounded. (It will be the case that the views of $S_0$ and $S_1$ together reveal $i$, but that is OK. Our assumption is that of non-collusion, that is $S_0$ and $S_1$ do not ever collude and put their views together.)

(b) *(4 points)* Unfortunately, the above protocol has a total communication complexity of $2(n + 1)$ bits, because the user $U$ sends an $n$-bit message and receives a one-bit response from each of the two servers. Note that there is a trivial protocol with communication complexity $n$ that is also private: one of the servers simply sends the entire string $x$ to the user $U$. We would like to reduce the amount of communication to significantly below $n$. Alter the scheme from part 2a in such a way as to reduce the amount of communication to $o(n)$.

(Hint: It might be useful to reorganise the database into a $\sqrt{n} \times \sqrt{n}$ table.)

(c) *(5 points)* Unfortunately, in the above protocol each server needs to look at the entire database to answer every query. That's too computationally expensive.

Your goal is to alter the scheme from 2b in such a way as to reduce the computation *per query* to $O(n/\log n)$, while keeping the communication complexity the same. In order to do that, you are allowed to preprocess the database into a polynomially larger string in an *offline* phase, that is before receiving any queries.

This preprocessing could be computationally expensive, that is, it could take polynomial time in $n$. However, the computation in the *online* phase, after receiving a query from the client, should take time $o(n)$. In particular, a solution that achieves online computational complexity $O(n/\log n)$ will receive full points.

3. *(6 points)* **Graphic Secret Sharing.**

In this problem, we consider interesting variants of secret-sharing.

(a) *(3 points)* (Path-out-of-graph secret sharing.) You have been approached by the king of Pessiland, who wishes to establish a series of qualifying trials by which his squires may be elevated to the knighthood. He has built a labyrinth of bridges and islands in the moat around his castle such that each bridge is manned by a guard. The king wants to design the maze such that, in order to succeed, the squire being tried must recover a secret piece of information from each guard that he passes on the way from the starting island (node $s$) to the king's castle (node $t$), such that all the secrets that the squire gathers along his journey can be put together after he reaches the king's castle into a password which will cause the guards at the castle to unlock the front gate for him. However, the king wishes the squire to be able to succeed regardless of which path he takes through the maze from the starting node $s$ to the finishing node $t$.

Fix some prime $p$, and consider some secret $x$ which is chosen uniformly at random from $\mathbb{Z}_p$. Model the labyrinth as an undirected, connected graph $G = (V, E)$ where $|V| = n$ and $|E| = m$. Suppose there is a value in $\mathbb{Z}_p$ associated with each distinct edge in the graph, which is the secret held by the guard responsible for that edge. Fix two distinct vertices $s, t \in V$ corresponding to the starting node $s$ and the ending vertex $t$.

Now, the king wants to assign values in $\mathbb{Z}_p$ to each of the $m$ distinct edges so that the following properties hold.

- For any subset $E' \subseteq E$ so that the corresponding edges contain a path from vertex $s$ to vertex $t$, the values assigned to the edges $e \in E'$ can be put together to recover the secret $x$.

- For any subset $E' \subseteq E$ so that the corresponding edges contain no path from vertex $s$ and $t$ (i.e. the corresponding subgraph contains $s$ and $t$ in different components, or not at all), the values assigned to all parties $m \in M'$ information-theoretically reveal nothing about $x$. Formally speaking, the distribution of $x$ conditioned on the values of all the shares assigned to parties $m \in M'$ should remain uniform.

Describe an assignment protocol that the king can run such that the above properties hold, and prove that they hold.

(Hint: It might be useful to associate each vertex with a random value in $\mathbb{Z}_p$.)

(b) *(3e points)* (Democracy of democracies) The king of Pessiland has organised his barons and knights so that each of the $2k+1$ barons have $2k+1$ loyal knights that only answer to their respective barons, where $k \geq 1$. One day, the king decides that he wants to go on an adventure across the seas to find Cryptomania—a magical land which is said to keep many secrets that the king longs to learn.

Completing the journey to Cryptomania requires an adventurer from Pessiland to surmount many dangerous obstacles, such as crossing the tempestuous Channel of Publishing in order to reach the comparative safety of the Tenure Isles, before traversing the warm tropical forest in which the Plague of Crackpottery and False Hope is endemic. Since everybody else who has attempted to find Cryptomania has expired in the attempt, the king decides to leave a will in his safe, which can be unlocked with a code $s$ that he chooses uniformly at random from $\mathbb{Z}_p$. However, he only wants the code to his safe to be revealed if a sufficient number of his loyal barons and knights decide that he is indeed gone never to return. In particular, the king wants to assign values in $\mathbb{Z}_p$ to each of his barons and knights so that the following properties hold for every set $B' \subseteq B$, where $B$ is the set of all the king's barons and knights,

- If for at least $k+1$ of the barons in $B'$, $k+1$ of their knights are also in $B'$, they can together recover the code $s$.
- Otherwise, the parties in $B'$ information-theoretically learn nothing about the code $s$. Formally speaking, the distribution of $s$ conditioned on the values of all the shares assigned to parties $b \in B'$ should remain uniform.

Describe an assignment protocol that the king can use such that the above properties hold, and prove that they hold.

4. *(8 points)* **Subset-revealing encryption.**

Professor Moriarty runs a crime ring in 1891, and a large number of criminals report to him. He would like to communicate with all of his subordinates by posting encrypted bulletins in the agony column of *The Times*. However, he would like to keep his subordinates from knowing about operations in which they have no part, in order to minimise the risk of information leakage if any one of them gets caught and confesses or rats out his cronies. For this purpose, he asks you, a time traveller from 2022, to design a *subset-revealing encryption scheme* consisting of the following algorithms:

- MasterGen$(1^\lambda)$ generates a secret master key $msk$.
- Enc$(msk, m)$ takes in a message $m \in \{0,1\}^n$ and produces a ciphertext $c$. (For simplicity, we will assume that $m$ is always of length $n(\lambda)$ for some large polynomial $n(\cdot)$.)
- KeyGen$(msk, T)$ takes in the master key $msk$ and a subset $T \subseteq [n]$, and outputs a secret key $sk_T$. We require that $sk_T$ is *succinct*, in that its length is always $O(\lambda)$ regardless of the size of $T$.

- DecMaster($msk, c$) outputs a message $m$. We require that Enc and DecMaster satisfy perfect correctness.

- Dec($sk_T, c$) outputs $\{m_i\}_{i \in T}$, where $m = $ DecMaster($msk, c$).

We say that this scheme satisfies *subset-revealing correctness* if, for any $T$ and any $m \in \{0,1\}^n$, Dec($sk_T$, Enc($msk, m$)) $= m_T$, where $m_T$ denotes the restriction of $m = $ DecMaster($msk$, Enc($msk, m$)) to the coordinates in $T$.

The security requirement for this scheme is as follows. Consider the following security game between a challenger $\mathcal{C}$ and a PPT adversary $\mathcal{A}$:

1. The challenger generates $msk \leftarrow$ Gen($1^\lambda$).

2. $\mathcal{A}$ generates a subset $T \subseteq [n]$ and sends $T$ to the challenger. The challenger responds by providing $sk_T \leftarrow$ KeyGen($msk, T$) to $\mathcal{A}$.

3. $\mathcal{A}$ submits any message $m_1 \in \{0,1\}^n$ to the challenger, and the challenger responds with Enc($msk, m_1$).

4. Step 3 is repeated an arbitrary polynomial in $\lambda$ many times.

5. $\mathcal{A}$ outputs two messages $m, m'$ such that $m_T = m'_T$, where where $m_T$ denotes the restriction of a message $m$ to the coordinates in $T$.

6. The challenger picks $m^* \leftarrow \{m, m'\}$ uniformly at random and sends $\mathcal{A}$ the encryption Enc($msk, m^*$).

We say the scheme is *non-adaptively subset-revealing secure* if no PPT $\mathcal{A}$ can guess which message the challenger chose in step 6 with better than $\frac{1}{2} + $ negl($\lambda$) probability.

To use this encryption scheme, Moriarty, who holds the master key, can encrypt any message using the master key and send the resulting ciphertext to the agony column of *The Times*. At the same time that he generates his master key, Moriarty can also produce secret keys corresponding to different subsets $T$ for each of his subordinates, in accordance with the operations that they are part of.

Let $\mathbb{G}$ be a group of prime order $p(\lambda) = 2^{O(\lambda)}$, and let $g$ be a generator of $\mathbb{G}$. Consider a subset-revealing encryption scheme whose MasterGen, Enc, DecMaster algorithms are defined as follows:

- MasterGen($1^\lambda$) samples $y_1, \ldots, y_n \leftarrow \mathbb{Z}_p$ independently and uniformly at random and outputs these.

- Enc($msk, m$) samples $r_1, \ldots, r_n$ independently and uniformly at random from $\mathbb{Z}_p$ and outputs $g^{r_1}, \ldots, g^{r_n}$ in addition to the $n \times n$ matrix $E$ whose $(i, j)$th entry is

$$E_{ij} = \begin{cases} g^{r_i y_j + m_i} & i = j \\ g^{r_i y_j} & \text{else} \end{cases}.$$

- DecMaster$(msk, c)$ computes the matrix $M$ whose $(i, j)$th entry is $g^{r_i y_j}$ using $msk$ and $g^{r_1}, \ldots, g^{r_n}$, and for all $i$ computes $E_{ii} \cdot M_{ii}^{-1} = \begin{cases} g & m_i = 1 \\ 1 & m_i = 0 \end{cases}$.

Define the algorithms KeyGen and Dec for this encryption scheme such that the scheme satisfies subset-revealing correctness, and prove its non-adaptive subset-revealing security assuming the DDH assumption in $\mathbb{G}$. Your KeyGen algorithm should be such that KeyGen$(msk, T)$ always outputs a *single element* of $\mathbb{Z}_p$.

Hint: You may cite past homework problems, but we recommend that you do the security reductions in this problem carefully. A good sanity check is whether or not you formally used all the relevant assumptions (e.g.: where do you use that $m_T = m'_T$? where do you use that $y_1, \ldots, y_n$ are sampled independently?). You should also check carefully that your reduction's answers to all of $\mathcal{A}$'s queries are distributed how $\mathcal{A}$ expects them to be distributed. If you handwave, Professor Moriarty will not be happy, and he is 'not a man who lets the grass grow under his feet'.

*Remark.* The definition of security which we gave is non-adaptive in the sense that $\mathcal{A}$ is not allowed to submit encryption queries before it provides the subset $T$ to the challenger. Can you come up with an *adaptively* secure scheme, in which $\mathcal{A}$ is allowed to submit encryption queries both before and after it provides the subset $T$ to the challenger?

---

**Solution**

Define $sk_T = \sum_{i \in T} y_i$. Subset-revealing correctness can be verified by direct calculation.

---